Lab 2

Divya K Raman, EE15B085

Answer 1

CAP Theorem is comprised of three technical terms. C stands for Consistency where all nodes see the data in homogeneous form i.e. every node has the same knowledge of data at any instant of time. A stands for availability i.e. A guarantee that every request receives a response which may be processed or failed. P stands for Partition Tolerance i.e. the system continues to operate even if a message is lost or part of the system fails Any system which satisfies two out of these three properties is a distributed system. Hadoop supports the Availability and Partition Tolerance property. The Consistency property is not supported because only namenode has the information of where the replicas are placed. This information is not available with each and every node of the cluster.

Answer 2

Apache Hive provides a database query interface which resembles SQL to data stored in databases and file systems that integrate with Hadoop. Pig enables data workers to write complex data transformations and manipulations without knowing Java. Both Pig and Hive are feature complete. Either of them can be used to do a particular task. Depending on their primary use case, data scientists generally prefer one over the other. Hive is easier for users proficient with SQL to query data whereas Pig has data flow strengths. Pig can bring data into Hadoop and get it into the form suitable for querying.

Answer 3

Before uploading any of the 4 files:

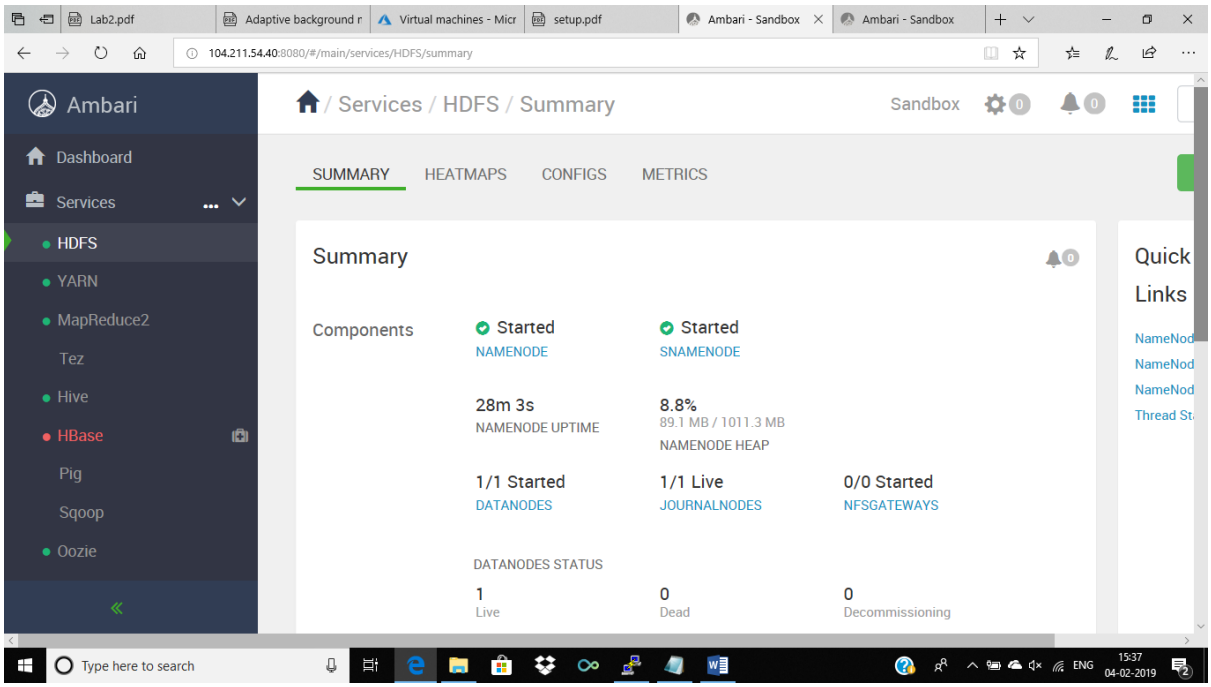

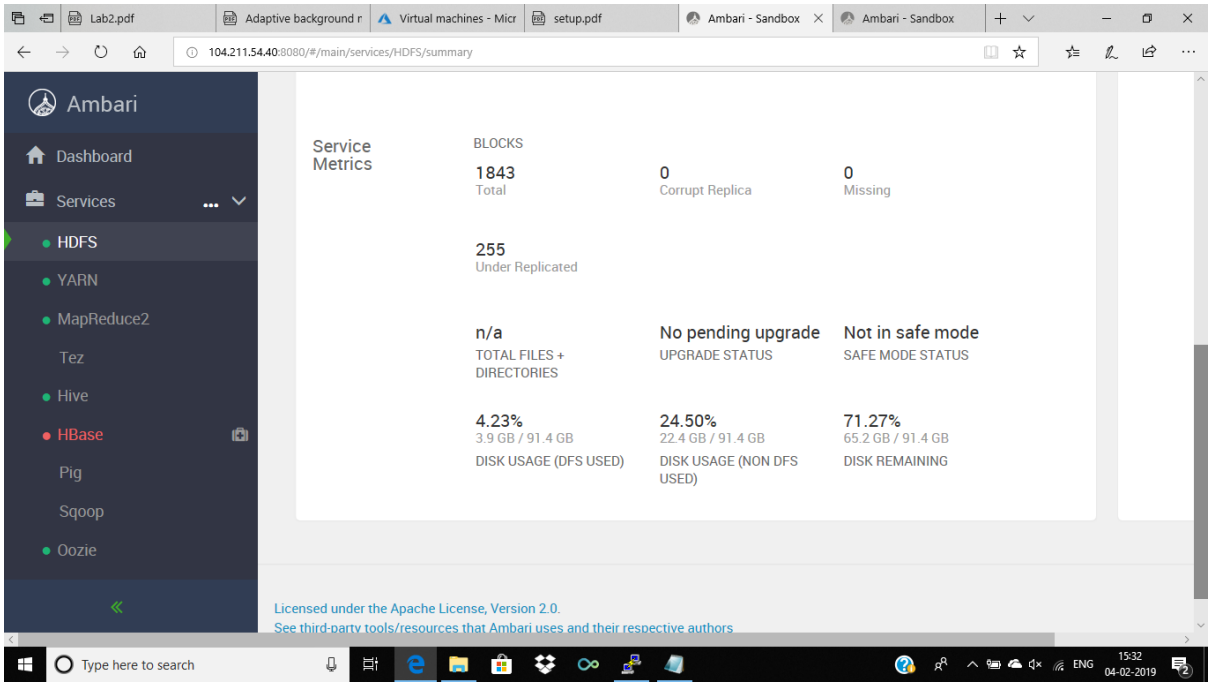

After uploading trucks.csv

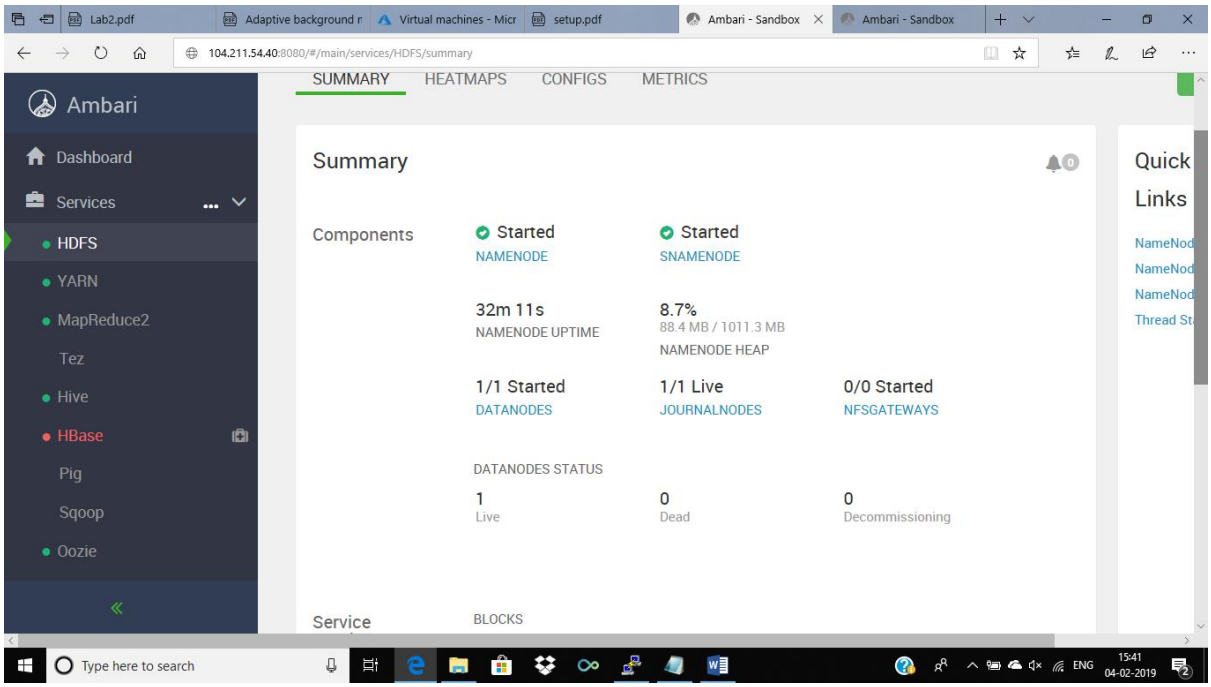

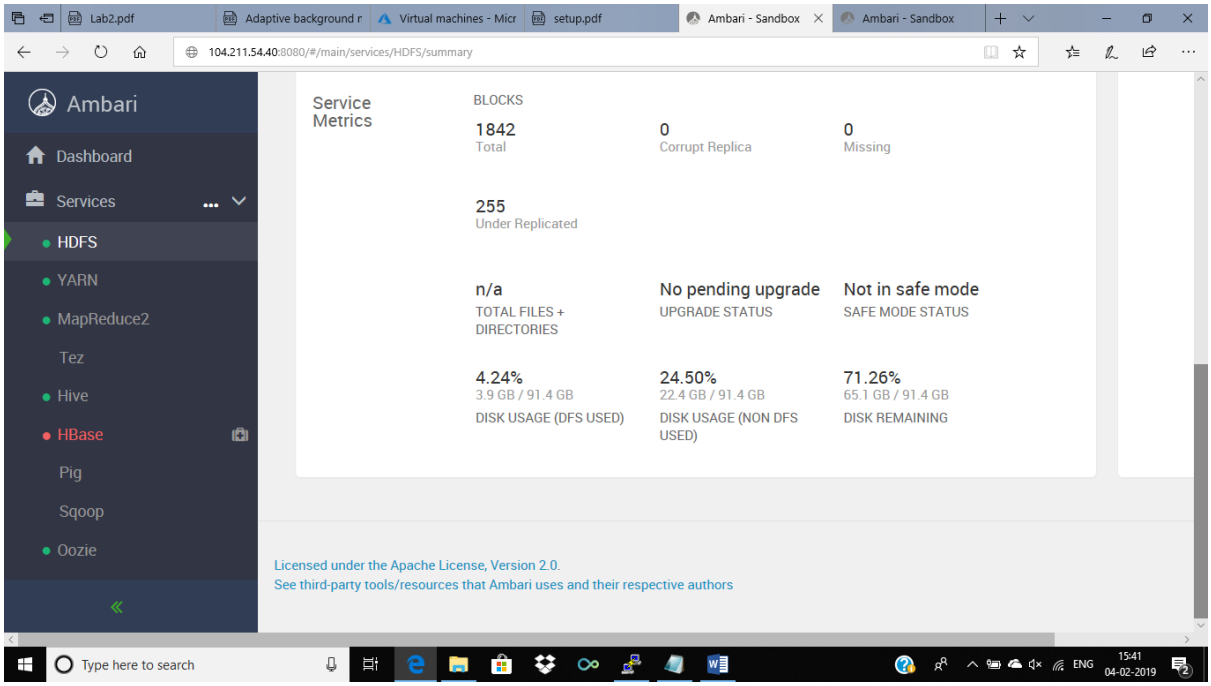

After uploading seasons.csv

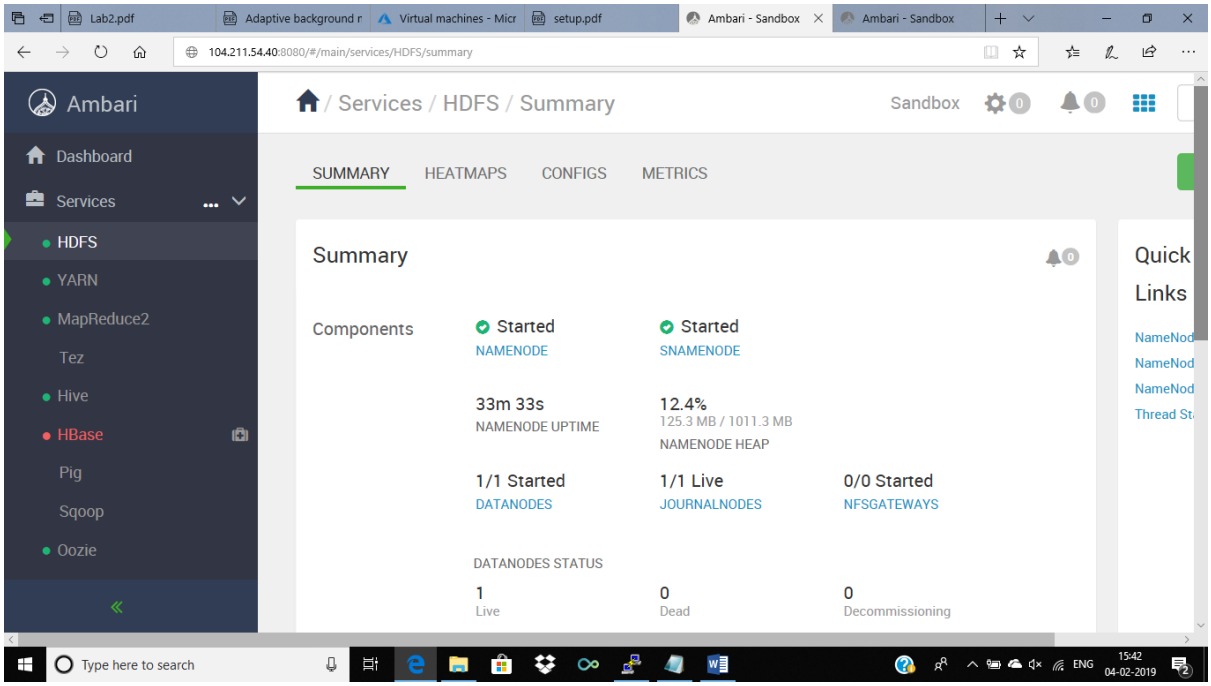

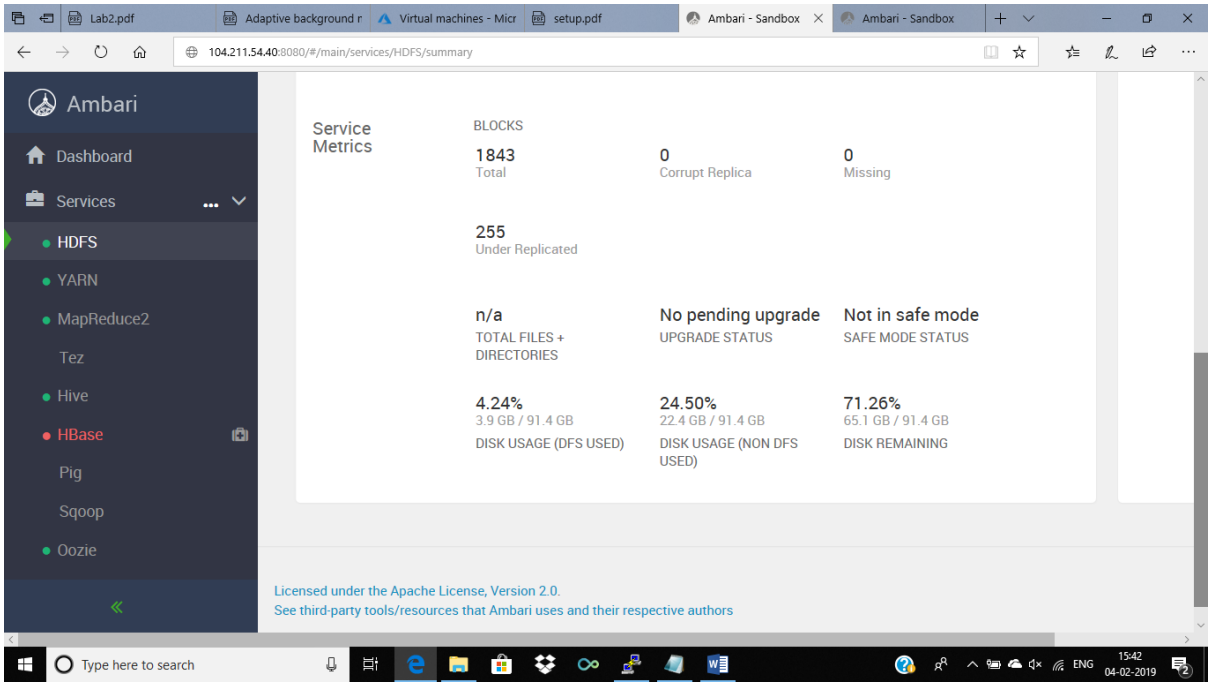

After products.tsv

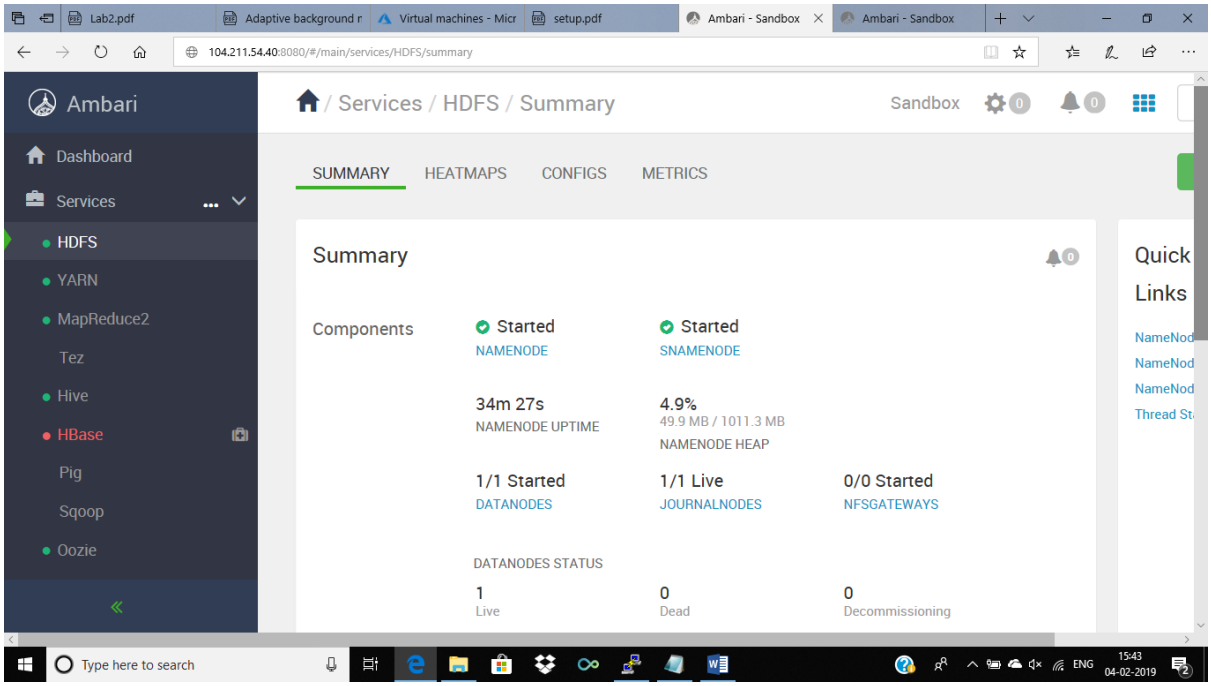

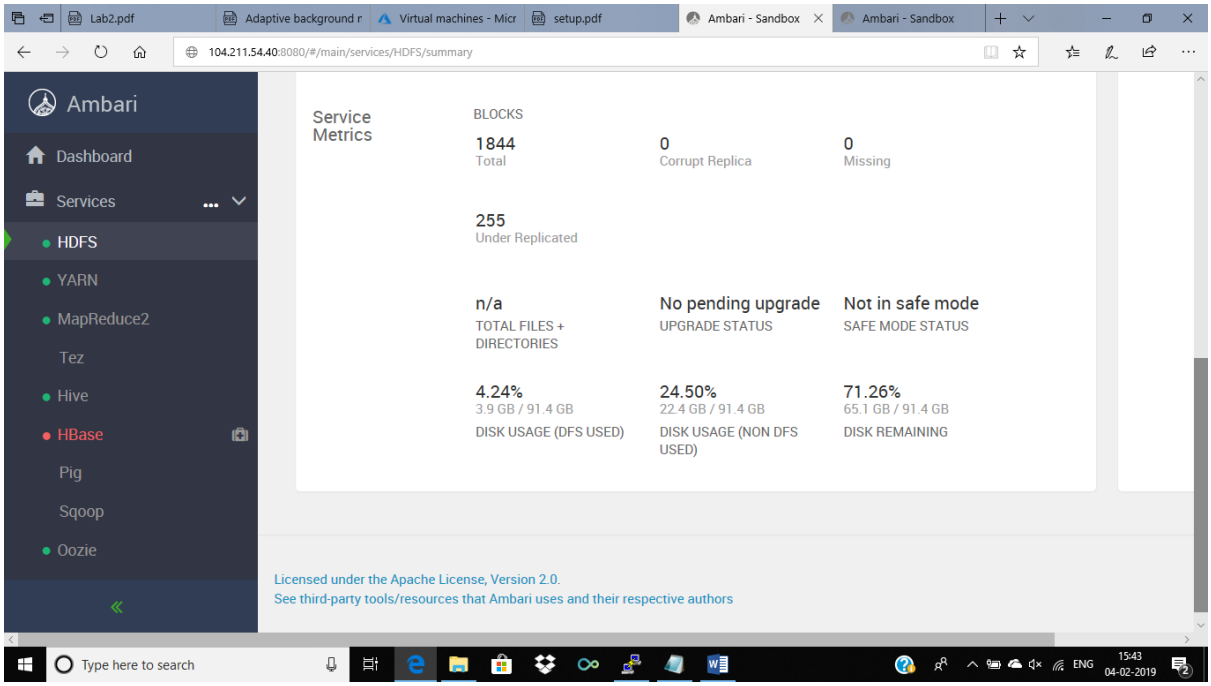

As files are added to the server, number of blocks increases which is understandable as new files are being added to the server. Under replicated remains the same. Under replicated blocks are blocks that do not meet their target replication for the file they belong to. This number indicates the number of copies of blocks missing from the environment. While uploading files, there were probably no copies of blocks which went missing and hence this number remains unchanged. Heap size first increases and then decreases. In Java, heap memory is an area of memory reserved for data that is created at runtime that is, when the program actually executes. NameNode heap size depends on many factors, such as the number of files, the number of blocks, and the load on the system.

Answer 4

In map, each worker node applies the map function to the local data, and writes the output to a temporary storage. A master node ensures that only one copy of redundant input data is processed. In Reduce, worker nodes now process each group of output data, per key, in parallel. In the given problem, each row describes a match officiated by a referee. In the map stage, we map the matches to the corresponding referees. In the reduce stage, we combine the matches officiated by each referee and take a total count for each referee.

Answer 5

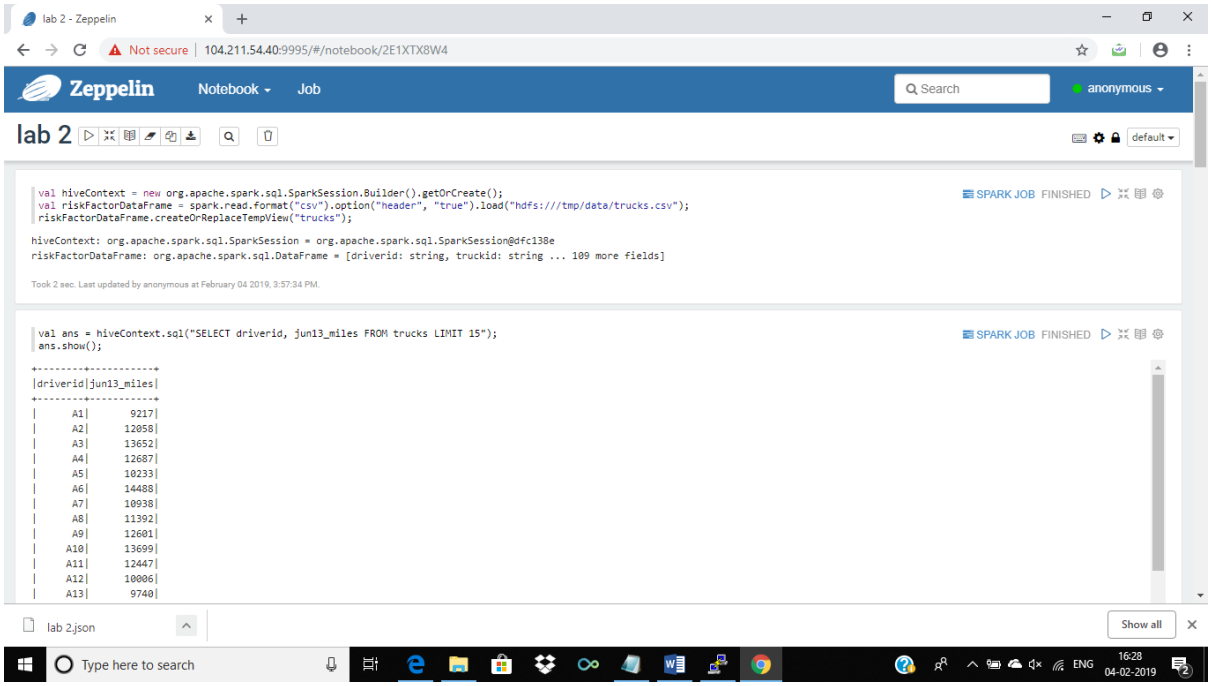
lab 2 - Zeppelin
Not secure | 104.211.54.40:9995/#/notebook/2E1XTX8W4
Zeppelin
Notebook
Job
Search
anonymous
lab 2
default
val hiveContext = new org.apache.spark.sql.SparkSession.Builder().getOrCreate();
val riskFactorDataFrame = spark.read.format("csv").option("header", "true").load("hdfs:///tmp/data/trucks.csv");
riskFactorDataFrame.createOrReplaceTempView("trucks");
SPARK JOB FINISHED
hiveContext: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@dfc138e
riskFactorDataFrame: org.apache.spark.sql.DataFrame = [driverid: string, truckid: string ... 109 more fields]
Took 2 sec. Last updated by anonymous at February 04 2019, 3:57:34 PM.
val ans = hiveContext.sql("SELECT driverid, jun13_miles FROM trucks LIMIT 15");
ans.show();
SPARK JOB FINISHED
+--------+-----------+
|driverid|jun13_miles|
+--------+-----------+
A1	9217
A2	12058
A3	13652
A4	12687
A5	10233
A6	14488
A7	10938
A8	11392
A9	12601
A10	13699
A11	12447
A12	10006
A13	9740
lab 2.json
Show all
Type here to search
ENG
16:28
04-02-2019

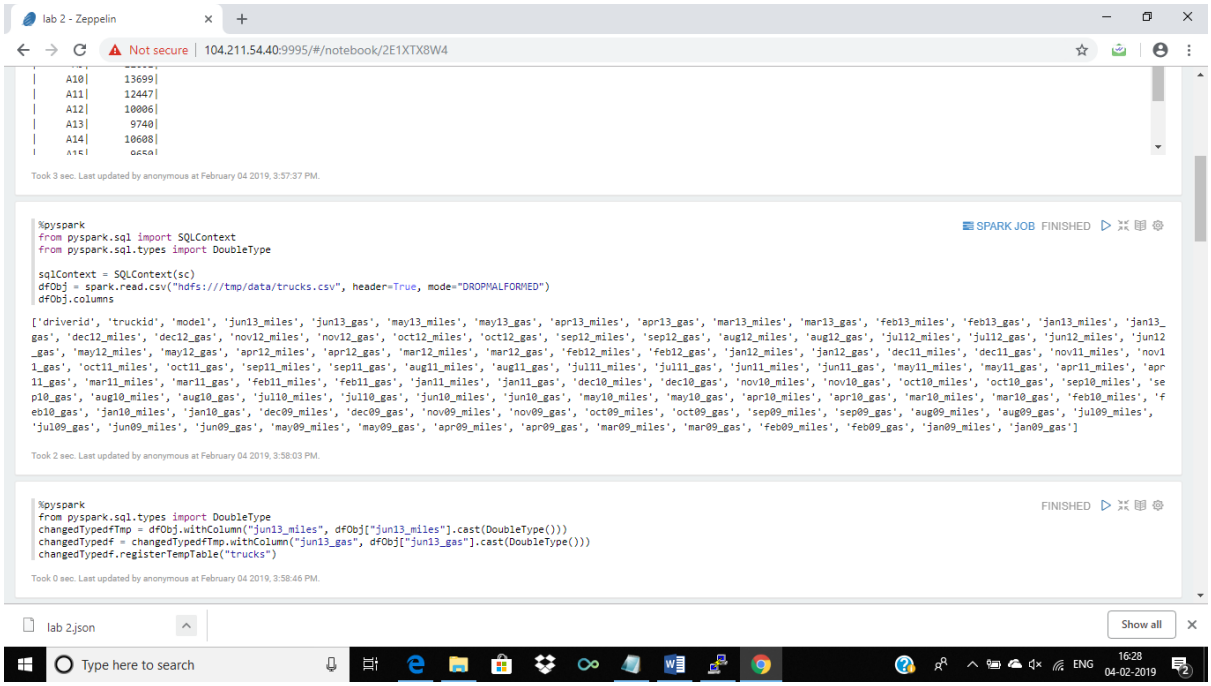
lab 2 - Zeppelin
Not secure | 104.211.54.40:9995/#/notebook/2E1XTX8W4
A10	13699
A11	12447
A12	10006
A13	9740
A14	10608
Took 3 sec. Last updated by anonymous at February 04 2019, 3:57:37 PM.
%pyspark
from pyspark.sql import SQLContext
from pyspark.sql.types import DoubleType
sqlContext = SQLContext(sc)
dfObj = spark.read.csv("hdfs:///tmp/data/trucks.csv", header=True, mode="DROPMALFORMED")
dfObj.columns
SPARK JOB FINISHED
['driverid', 'truckid', 'model', 'jun13_miles', 'jun13_gas', 'may13_miles', 'may13_gas', 'apr13_miles', 'apr13_gas', 'mar13_miles', 'mar13_gas', 'feb13_miles', 'feb13_gas', 'jan13_miles', 'jan13_gas', 'dec12_miles', 'dec12_gas', 'nov12_miles', 'nov12_gas', 'oct12_miles', 'oct12_gas', 'sep12_miles', 'sep12_gas', 'aug12_miles', 'aug12_gas', 'jul12_miles', 'jul12_gas', 'jun12_miles', 'jun12_gas', 'may12_miles', 'may12_gas', 'apr12_miles', 'apr12_gas', 'mar12_miles', 'mar12_gas', 'feb12_miles', 'feb12_gas', 'jan12_miles', 'jan12_gas', 'dec11_miles', 'dec11_gas', 'nov11_miles', 'nov11_gas', 'oct11_miles', 'oct11_gas', 'sep11_miles', 'sep11_gas', 'aug11_miles', 'aug11_gas', 'jul11_miles', 'jul11_gas', 'jun11_miles', 'jun11_gas', 'may11_miles', 'may11_gas', 'apr11_miles', 'apr11_gas', 'mar11_miles', 'mar11_gas', 'feb11_miles', 'feb11_gas', 'jan11_miles', 'jan11_gas', 'dec10_miles', 'dec10_gas', 'nov10_miles', 'nov10_gas', 'oct10_miles', 'oct10_gas', 'sep10_miles', 'sep10_gas', 'aug10_miles', 'aug10_gas', 'jul10_miles', 'jul10_gas', 'jun10_miles', 'jun10_gas', 'may10_miles', 'may10_gas', 'apr10_miles', 'apr10_gas', 'mar10_miles', 'mar10_gas', 'feb10_miles', 'feb10_gas', 'jan10_miles', 'jan10_gas', 'dec09_miles', 'dec09_gas', 'nov09_miles', 'nov09_gas', 'oct09_miles', 'oct09_gas', 'sep09_miles', 'sep09_gas', 'aug09_miles', 'aug09_gas', 'jul09_miles', 'jul09_gas', 'jun09_miles', 'jun09_gas', 'may09_miles', 'may09_gas', 'apr09_miles', 'apr09_gas', 'mar09_miles', 'mar09_gas', 'feb09_miles', 'feb09_gas', 'jan09_miles', 'jan09_gas']
Took 2 sec. Last updated by anonymous at February 04 2019, 3:58:03 PM.
%pyspark
from pyspark.sql.types import DoubleType
changedTypedfTmp = dfObj.withColumn("jun13_miles", dfObj["jun13_miles"].cast(DoubleType()))
changedTypedf = changedTypedfTmp.withColumn("jun13_gas", dfObj["jun13_gas"].cast(DoubleType()))
changedTypedf.registerTempTable("trucks")
FINISHED
Took 0 sec. Last updated by anonymous at February 04 2019, 3:58:46 PM.
lab 2.json
Show all
Type here to search
ENG
16:28
04-02-2019

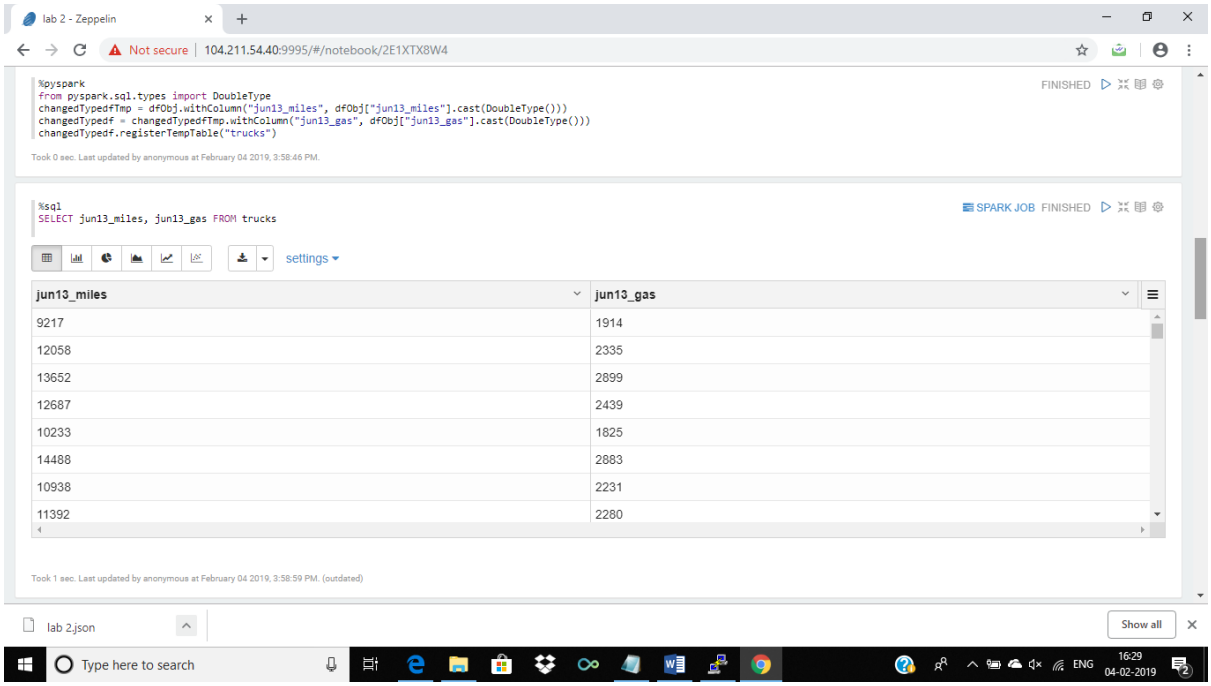
lab 2 - Zeppelin
Not secure | 104.211.54.40:9995/#/notebook/2E1XTX8W4
%pyspark
from pyspark.sql.types import DoubleType
changedTypedfTmp = dfObj.withColumn("jun13_miles", dfObj["jun13_miles"].cast(DoubleType()))
changedTypedf = changedTypedfTmp.withColumn("jun13_gas", dfObj["jun13_gas"].cast(DoubleType()))
changedTypedf.registerTempTable("trucks")
FINISHED
Took 0 sec. Last updated by anonymous at February 04 2019, 3:58:46 PM.
%sql
SELECT jun13_miles, jun13_gas FROM trucks
SPARK JOB FINISHED
settings
jun13_miles
jun13_gas
9217 1914
12058 2335
13652 2899
12687 2439
10233 1825
14488 2883
10938 2231
11392 2280
Took 1 sec. Last updated by anonymous at February 04 2019, 3:58:59 PM. (outdated)
lab 2.json
Show all
Type here to search
ENG
16:29
04-02-2019

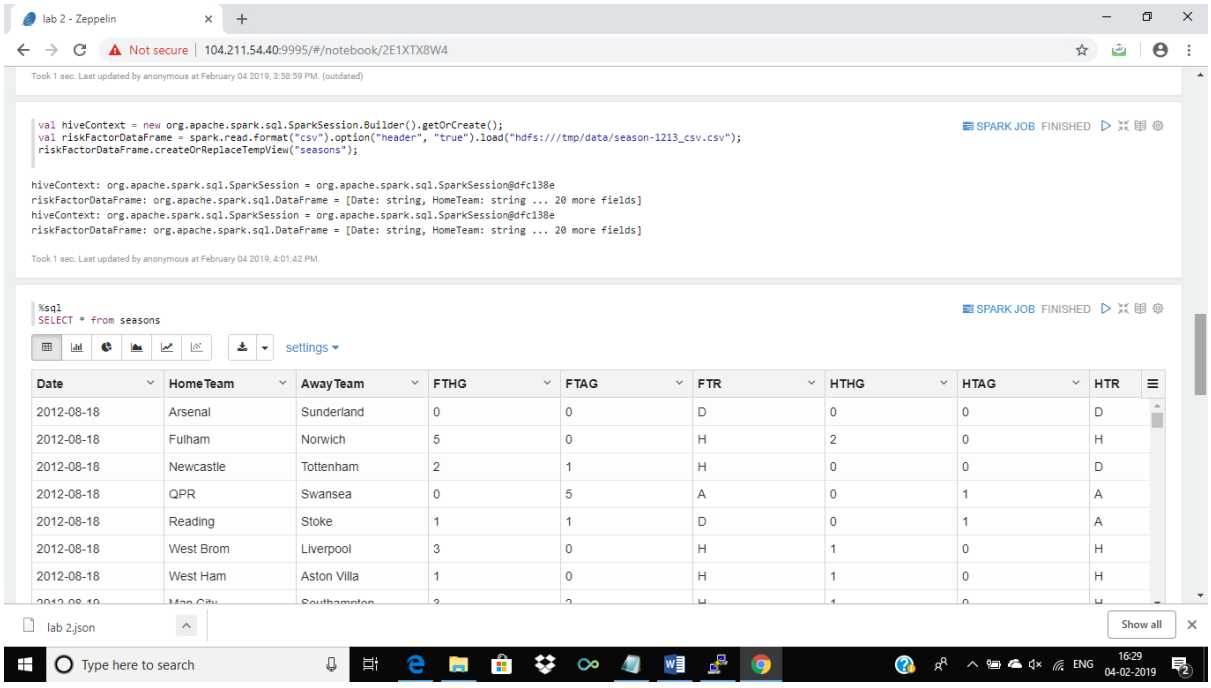
lab 2 - Zeppelin
Not secure | 104.211.54.40:9995/#/notebook/2E1XTX8W4
Took 1 sec. Last updated by anonymous at February 04 2019, 3:58:59 PM. (outdated)
val hiveContext = new org.apache.spark.sql.SparkSession.Builder().getOrCreate();
val riskFactorDataFrame = spark.read.format("csv").option("header", "true").load("hdfs:///tmp/data/season-1213_csv.csv");
riskFactorDataFrame.createOrReplaceTempView("seasons");
SPARK JOB FINISHED
hiveContext: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@dfc138e
riskFactorDataFrame: org.apache.spark.sql.DataFrame = [Date: string, HomeTeam: string ... 20 more fields]
hiveContext: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@dfc138e
riskFactorDataFrame: org.apache.spark.sql.DataFrame = [Date: string, HomeTeam: string ... 20 more fields]
Took 1 sec. Last updated by anonymous at February 04 2019, 4:01:42 PM.
%sql
SELECT * from seasons
SPARK JOB FINISHED
settings
Date HomeTeam AwayTeam FTHG FTAG FTR HTHG HTAG HTR
2012-08-18 Arsenal Sunderland 0 0 D 0 0 D
2012-08-18 Fulham Norwich 5 0 H 2 0 H
2012-08-18 Newcastle Tottenham 2 1 H 0 0 D
2012-08-18 QPR Swansea 0 5 A 0 1 A
2012-08-18 Reading Stoke 1 1 D 0 1 A
2012-08-18 West Brom Liverpool 3 0 H 1 0 H
2012-08-18 West Ham Aston Villa 1 0 H 1 0 H
lab 2.json
Show all
Type here to search
ENG
16:29
04-02-2019

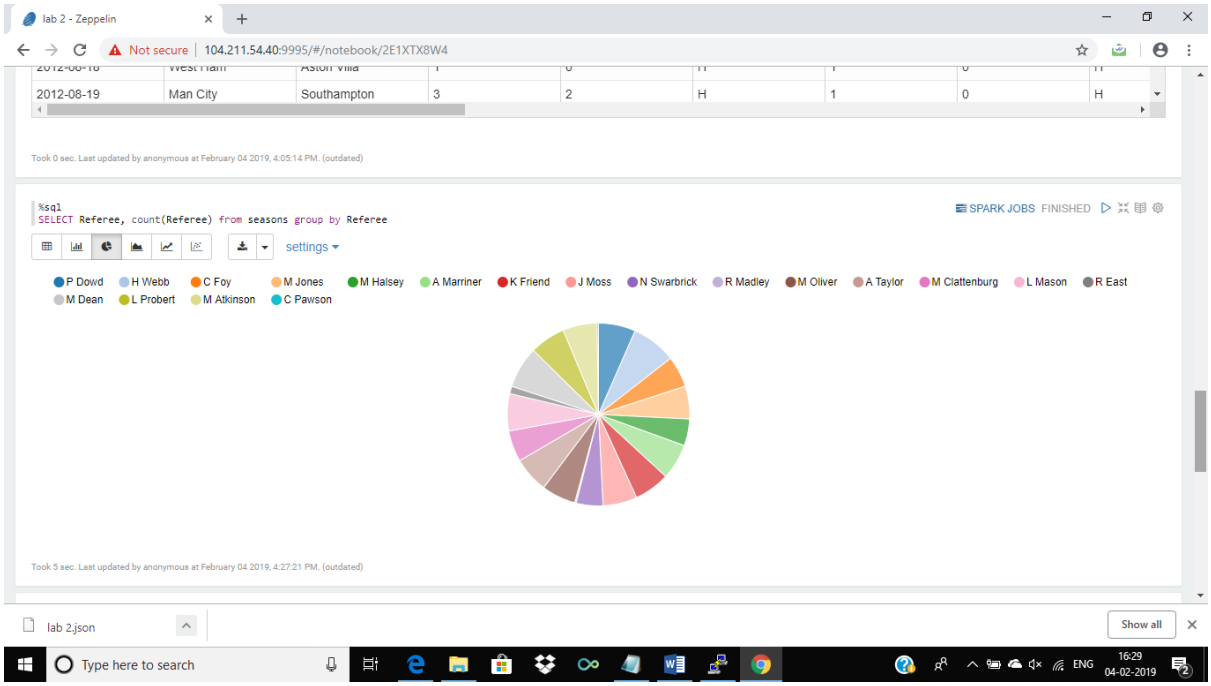
lab 2 - Zeppelin
Not secure | 104.211.54.40:9995/#/notebook/2E1XTX8W4
2012-08-19 Man City Southampton 3 2 H 1 0 H
Took 0 sec. Last updated by anonymous at February 04 2019, 4:05:14 PM. (outdated)
%sql
SELECT Referee, count(Referee) from seasons group by Referee
SPARK JOBS FINISHED
settings
P Dowd H Webb C Foy M Jones M Halsey A Marriner K Friend J Moss N Swarbrick R Madley M Oliver A Taylor M Clattenburg L Mason R East
M Dean L Probert M Atkinson C Pawson
Took 5 sec. Last updated by anonymous at February 04 2019, 4:27:21 PM. (outdated)
lab 2.json
Show all
Type here to search
ENG
16:29
04-02-2019

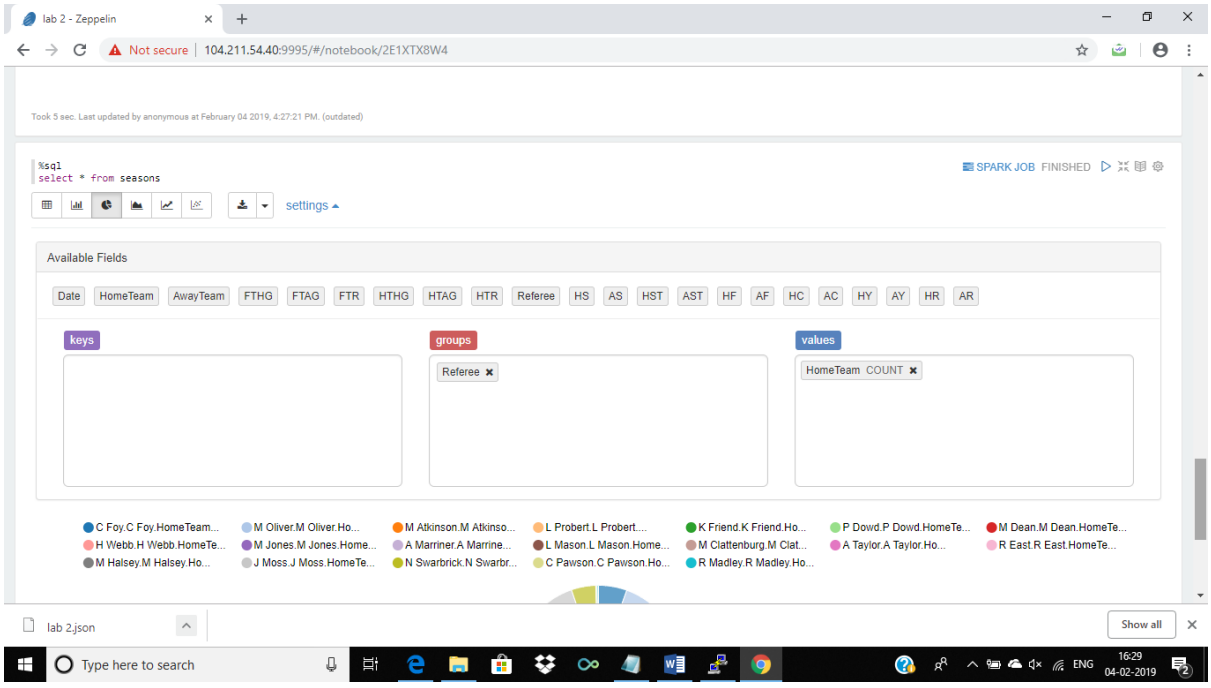
lab 2 - Zeppelin
Not secure | 104.211.54.40:9995/#/notebook/2E1XTX8W4
Took 5 sec. Last updated by anonymous at February 04 2019, 4:27:21 PM. (outdated)
%sql
select * from seasons
SPARK JOB FINISHED
settings
Available Fields
Date HomeTeam AwayTeam FTHG FTAG FTR HTHG HTAG HTR Referee HS AS HST AST HF AF HC AC HY AY HR AR
keys
groups
Referee
values
HomeTeam COUNT
C Foy.C Foy.HomeTeam... M Oliver.M Oliver.Ho... M Atkinson.M Atkinso... L Probert.L Probert.... K Friend.K Friend.Ho... P Dowd.P Dowd.HomeTe... M Dean.M Dean.HomeTe...
H Webb.H Webb.HomeTe... M Jones.M Jones.Home... A Marriner.A Marrine... L Mason.L Mason.Home... M Clattenburg.M Clat... A Taylor.A Taylor.Ho... R East.R East.HomeTe...
M Halsey.M Halsey.Ho... J Moss.J Moss.HomeTe... N Swarbrick.N Swarbr... C Pawson.C Pawson.Ho... R Madley.R Madley.Ho...
lab 2.json
Show all
Type here to search
ENG
16:29
04-02-2019

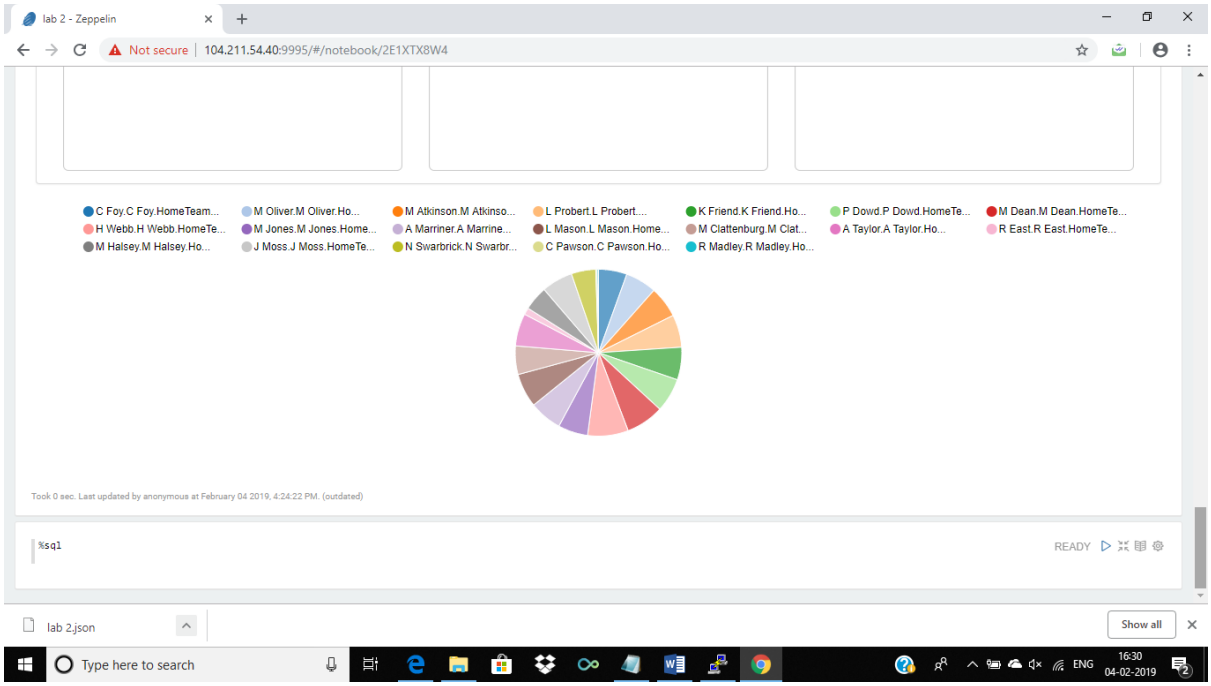

Number of blocks after running the notebooks:

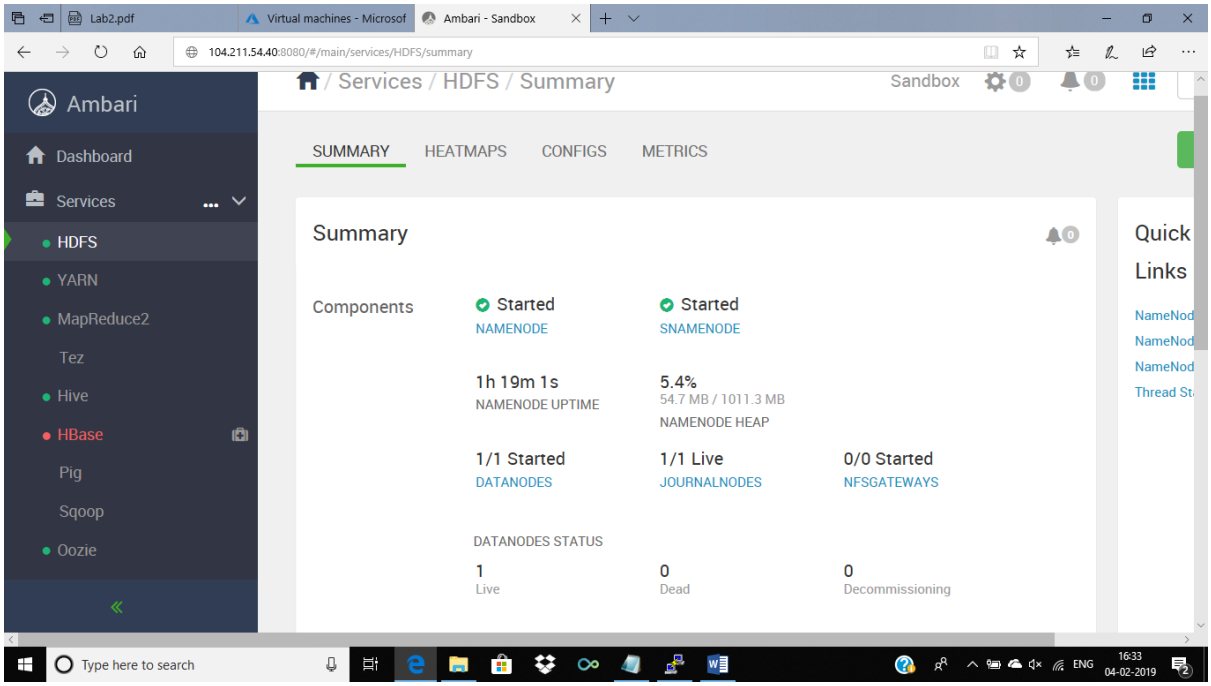

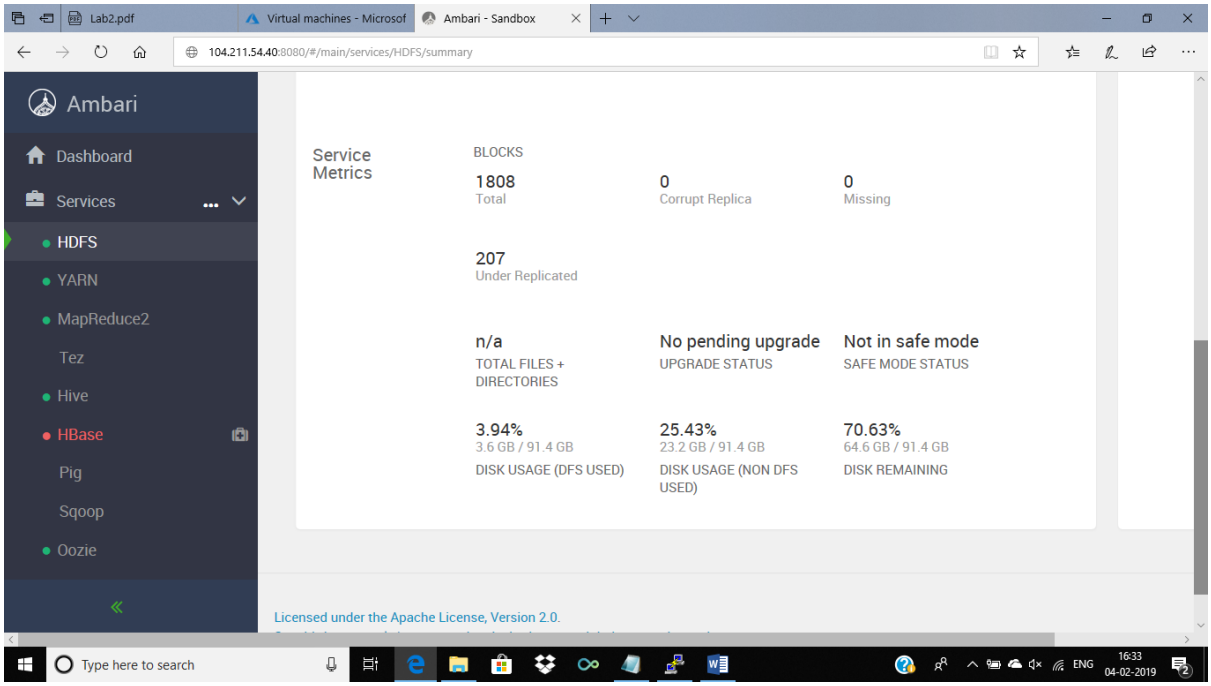

References:

https://data-flair.training/forums/topic/what-is-cap-theorem-what-aspects-hadoop-supports-from-this-theorem/

https://stackoverflow.com/questions/19923196/cap-with-distributed-system

https://hortonworks.com/tutorial/how-to-process-data-with-apache-hive/

https://hortonworks.com/tutorial/beginners-guide-to-apache-pig/

https://stackoverflow.com/questions/36977746/in-hadoop-whats-under-replication-and-over-replication-mean-and-how-does-it-work

https://docs.hortonworks.com/HDPDocuments/HDP2/HDP-2.6.5/bk_command-line-installation/content/configuring-namenode-heap-size.html